

# Class 19: Joins II

Last time

Relation S, N=500 pages  $P_S=80$   
 R, M=1000 pages  $P_R=100$

## Nested-Loop Joins

Simple  $(P_R \cdot M) \cdot N + M$  w/ R outer  
 Page-oriented  $M \cdot N + M$   
 Block-based  $\frac{M \cdot N}{K} + M$  w/ K buffer  
 Index  $M + M \cdot P_R \cdot (\text{index\_access\_cost} + \text{data\_access\_cost})$

hash  $\sim 1.2$   $\swarrow$   $\searrow$  B<sup>+</sup>-tree 2-4

$\swarrow$  clustered 1 I/O per page of matching tuples  
 $\searrow$  unclustered

## Sort-Merge Joins

$3 \cdot (M+N)$  if  $B > \sqrt{M}$  where M is # pages of the larger relation

$M+N$  if  $B > N$  where N corresponds to the smallest relation

1 I/O per matching tuple

## Today

- Hash Joins
- General Join Conditions
- Aggregates

## Hash Joins

- Use a hash function h to create partitions of both relations hashing (building)
- match tuples only between the corresponding partitions probing (matching)

B buffers




h hash function

$R \bowtie S$   
 $i=j$

building {  
 $\forall r \in R$   
read r and add it to buffer  $h(r_i)$   
 $\forall s \in R$   
read s and add it to buffer  $h(s_j)$




matching {  
for  $l=1, 2, \dots, K$   
 $\forall r \in R_l$   
read r and insert into in-memory AT using  $h_2(r_i)$   
 $\forall s \in S_l$   
read s and probe AT using  $h_2(s_j)$   
if match found add  $\langle r, s \rangle$  to the result  
clear hash table from memory to proceed  
with next pair of partitions

Building




Cost

$2 \cdot M$



$2 \cdot N$

## Matching



read every partition once  
in-memory HT w/  $h_2 (\neq h)$

Search in  $S_i$  as we stream  $R_i$

Cost:  $M + N$

total cost of Hash Join =  $3(M + N) = \boxed{4500} \rightarrow \boxed{93}$

## Memory Requirements

→ enough buffer for the largest partition of the smaller relation ( $S$ )

→ Input page for the other relation

→ Output page

→ a few pages of hash metadata

Fudge factor  $f$  (for example  $f = 1.04$ )

if  $h \rightarrow$  uniform

size of a partition  $\sim \frac{N}{B-1}$

$$B > \frac{f \cdot N}{B-1} + 2 \quad \Rightarrow \quad B > \sqrt{f \cdot N}$$

What if not enough memory? (for  $S_i$  to fit in memory)

→ apply the same algorithm recursively


→ read, repartition  $S_i, R_i$  with  $h_3 (\neq h_2, \neq h)$

→ matching per subpartition (mem. is enough)

\* if not, again recursion

What if we have more memory?

Hybrid  
Hash  
Join



Cost

→ hashing S  $N + N - \text{sizeof}(S_1)$   
 → hashing R  $M + M - \text{sizeof}(R_1)$   
 → matching  $M - \text{sizeof}(R_1) + N + \text{sizeof}(S_1)$

total  $3(M+N) - 2(\text{sizeof}(S_1) + \text{sizeof}(R_1))$

$B = 300$

$M = 1000$

$N = 500$

$$3(1000 + 500) - 2(500 + 250) = 4500 - 1500 = \boxed{3000}$$

$\boxed{6S}$

if  $B = 600$

read S once + build hash table

scan R once prob S on-the-fly

Hash Join vs

SMJ

cost  
mem. req

$$3(M+N)$$

$$B > \sqrt{R \cdot N} \leftarrow \text{smaller}$$

$$B > \sqrt{1.04 \cdot 500} = 23$$

$$3(M+N) - 2(\text{sizeof}(R_1) + \text{sizeof}(S_1))$$

$$M+N$$

$$3(M+N)$$

$$3(M+N)$$

$$B > \sqrt{M} \leftarrow \text{larger}$$

$$B > \sqrt{1000} = 32$$

$$3(M+N)$$

$$M+N$$

sorted

$$M+N$$

$$B(\geq m.r.) \leq N$$

$$B > N$$

output

if input sorted

BUT sensitive to data skew

(a) equality join on several attributes

(b) inequality joins

→ (a) for INLJ we need index with all attributes in join conditions

→ sort/hash use combination of all attributes

→ (b) INLJ w/ B<sup>+</sup>-Tree (not Hash Index)

HS/SMJ cannot work

Block NLJ the best approach

Set

UNION / EXCEPT (set difference)

→ Sorting

→ sort S+R on all attributes

→ merging → discard duplicates (UNION)  
→ set-difference

→ hashing

→ partition  $R + S$

→  $\nexists$  S-part probe corr. R-part

→ discard duplicates (UNION)  
→ set-difference

→ Intersection → special case of Join  
Equality across all attributes

---

## Aggregation

→ `SELECT AVG(sal) FROM E`

→ SCAN once

→ GROUP BY

$\langle \text{age, avg-salary} \rangle$

hash (age) →  $\langle \text{age, salary, count} \rangle$

sort (age) calculate "running info" of aggregation  
on-the-fly

→ if we have an index on  $\langle \text{Group-by, select, where} \rangle$   
can use only the index WAY FASTER

→ Buffering

Among things in parallel

tough to estimate what is needed by BP

SNLS BDN ✓.

BEN

LRU → sequential flooding.

MRU ✓.

INLS → sort the outer relation