

CS460: Intro to Database Systems

Class 21: Transactional Management Overview

Instructor: Manos Athanassoulis

<https://bu-disc.github.io/CS460/>

Administrativa – what lies ahead

WA5b (Bonus) – on indexing & LSM (deadline 11/17)

uploaded last week

WA6 – on Query Optimization (deadline 11/23)

coming on 11/17

PA2.1 – Query Optimization (deadline 11/30)

uploaded last week

PA2.2 – Row-stores vs. Column-stores (completion in the labs)

coming next week

WA7 (last WA) – on Transaction Processing (deadline 12/6)

coming next week

Final: in HAR 208 (this room) on Wednesday 12/15 3pm-5pm

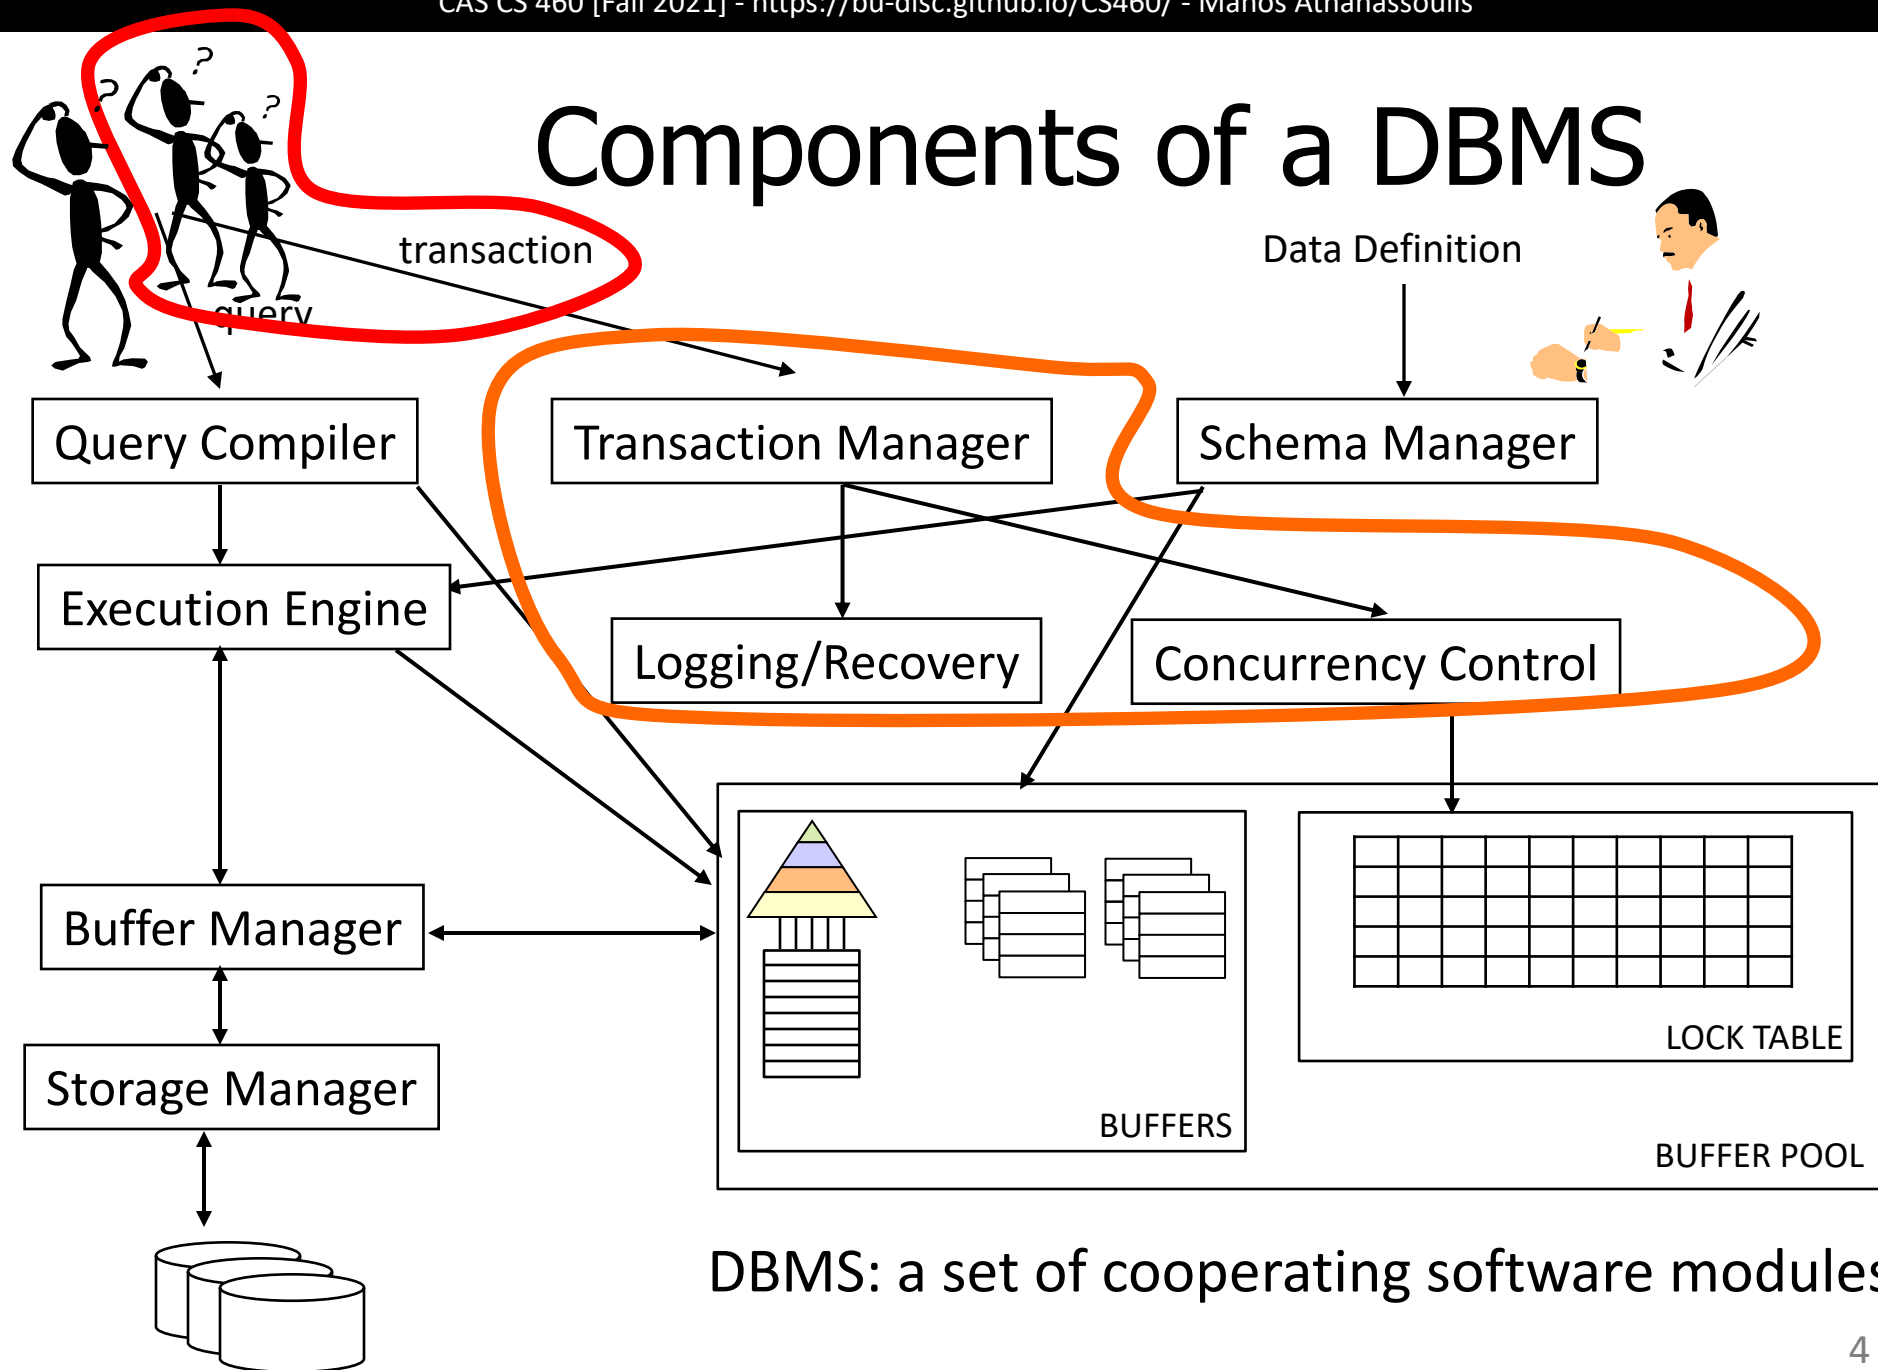
Transaction Management

Overview of ACID

Readings: Chapter 16.1

Concurrency control

Logging and recovery

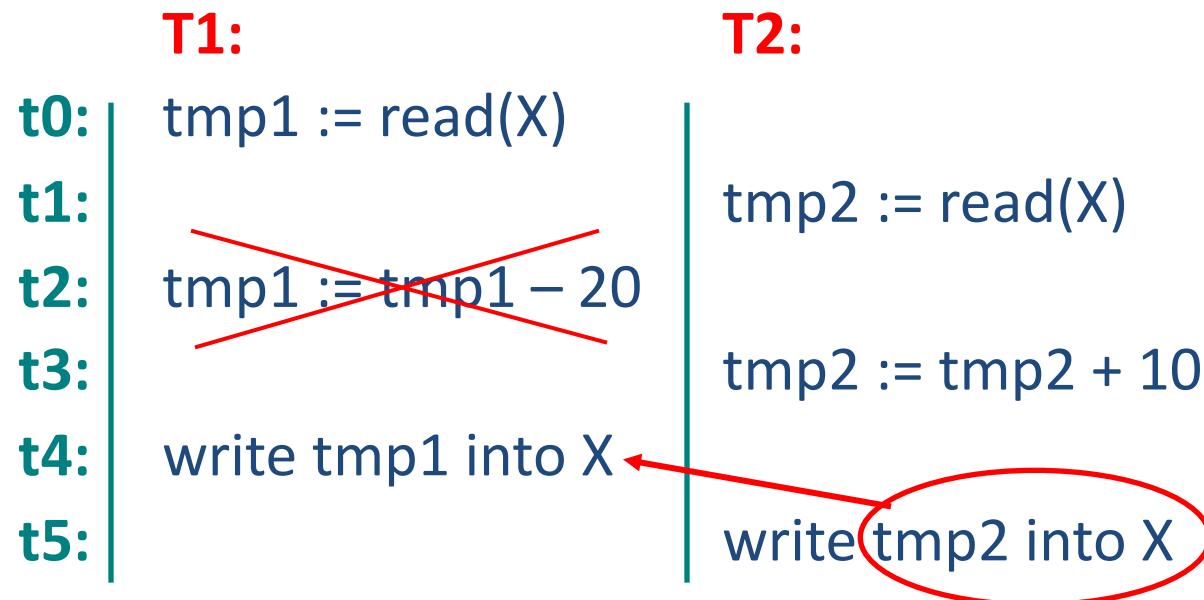


Problem Statement

Goal: concurrent execution of independent transactions

- utilization/throughput (“hide” waiting for I/Os)
- response time
- fairness

Example:



Arbitrary interleaving can lead to inconsistencies

Definitions

A program may carry out many operations on the data retrieved from the database

The DBMS is only concerned about what data is read/written from/to the database

database

a fixed set of named data objects (A, B, C, \dots)

transaction

a sequence of read and write operations ($read(A), write(B), \dots$)

Correctness: The **ACID** properties

A tomicity: All actions in the transaction happen, or none happen

C onsistency: If each transaction is consistent, and the DB starts consistent, it ends up consistent

I solation: Execution of one transaction is isolated from that of other transactions

D urability: If a transaction commits, its effects persist

Transaction Management

Overview of ACID

Concurrency control

Readings: Chapter 16.2-16.6


Logging and recovery

C Transaction Consistency

Consistency - data in DBMS is accurate in modeling real world and follows integrity constraints

User must ensure that transaction is consistent

Key point:



C Transaction Consistency (cont.)

Recall: Integrity constraints

- must be true for DB to be considered consistent
- **Examples:**
 1. FOREIGN KEY R.sid REFERENCES S
 2. ACCT-BAL ≥ 0

System checks integrity constraints and if they fail, the transaction rolls back (i.e., is aborted)

- Beyond this, DBMS does not understand data semantics
- e.g., how interest on a bank account is computed

I Isolation of Transactions

Users submit transactions, and

Each xact executes as if it was running **by itself**

- Concurrency is achieved by DBMS, which interleaves actions (reads/writes of DB objects) of various transactions.

Techniques for achieving isolation:

- Pessimistic – don't let problems arise in the first place
- Optimistic – assume conflicts are rare, deal with them *after* they happen.

I

Example

Consider two transactions:

| | | | | |
|-----|-------|-----------|----------|-----|
| T1: | BEGIN | A=A+100, | B=B-100 | END |
| T2: | BEGIN | A=1.06*A, | B=1.06*B | END |

1st xact transfers \$100 from B's account to A's

2nd xact credits both accounts with 6% interest

Assume at first A and B each have \$1000. What are the legal outcomes of running T1 and T2?

$$\$2000 * 1.06 = \$2120$$

There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together. **But, the net effect *must* be equivalent to these two transactions running serially in some order**

I

Example (Cont.)

Legal outcomes: $A=1166, B=954$ or $A=1160, B=960$

Remember: correct outcome: $A+B=\$2120$

Consider a possible interleaved schedule:

| | | |
|-----|-------------|------------|
| T1: | $A=A+100,$ | $B=B-100$ |
| T2: | $A=1.06*A,$ | $B=1.06*B$ |

This is OK (same as $T1;T2$). But what about:

| | | |
|-----|----------------------|-----------|
| T1: | $A=A+100,$ | $B=B-100$ |
| T2: | $A=1.06*A, B=1.06*B$ | |

Result: $A=1166, B=960; A+B = 2126$, bank loses \$6

The DBMS's view of the second schedule:

| | | |
|-----|--------------------------|--------------|
| T1: | $R(A), W(A),$ | $R(B), W(B)$ |
| T2: | $R(A), W(A), R(B), W(B)$ | |

I Anomalies with Interleaved Execution

Reading Uncommitted Data (WR Conflicts, “dirty reads”):

| | | |
|-----|---------------|-------------------|
| T1: | R(A), W(A), | R(B), W(B), Abort |
| T2: | R(A), W(A), C | |



Unrepeatable Reads (RW Conflicts):

| | | |
|-----|---------------|---------------|
| T1: | R(A), | R(A), W(A), C |
| T2: | R(A), W(A), C | |



I Anomalies (Continued)

Overwriting Uncommitted Data (WW Conflicts):

| | | |
|-----|-------|---------|
| T1: | W(A), | W(B), C |
| T2: | W(A), | W(B), C |



A gets its value from T2

B gets its values from T1

A correct execution would take both values from T2 or both from T1

I

Concurrency Control

How to avoid such anomalies?
“lock” data



Strict Two-phase Locking (Strict 2PL) Protocol

obtain an *S (shared) lock* on object before reading

obtain an *X (exclusive) lock* on object before writing

- (i) obtain locks automatically
- (ii) if a xact holds an X lock on object no other xact can acquire S or X
- (iii) if a xact holds an S lock, no other xact can acquire X (but only S)

2 phases: first acquire and then release all at the end
important: no lock is ever acquired after one has been released

Transaction Management

Overview of ACID

Concurrency control

Logging and recovery

Readings: Chapter 16.7

A Atomicity of Transactions



Two possible outcomes of executing a transaction:

- Transaction might *commit* after completing all its actions
- or it could *abort* (or be aborted by the DBMS) after executing some actions

DBMS guarantees that transactions are *atomic*.

- From user's point of view: transaction always either executes *all its actions*, or executes *no actions* at all

A Mechanisms for Ensuring Atomicity

One approach: **LOGGING**

- DBMS *logs* all actions so that it can *undo* the actions of aborted transactions

Another approach: **SHADOW PAGES**

- (ask me after class if you're curious)

Logging used by modern systems, because of the need for audit trail and for efficiency

Aborting a Transaction (i.e., Rollback)

If a xact T_i is aborted, all its actions must be undone

If T_j reads object last written by T_i , T_j must be aborted!

- Most systems avoid such *cascading aborts* by releasing locks only at end of the transaction (i.e., strict locking)
- If T_i writes an object, T_j can read it only after T_i finishes

To *undo* actions of an aborted transaction, DBMS maintains *log* which records every write

Log is also used to recover from system crashes:

- All active Xacts at time of crash are aborted when system comes back up



why?

to ensure atomicity!

The Log

Log consists of “records” that are written sequentially

- Typically chained together by transaction id
- Log is often *archived* on stable storage

Need for **UNDO** and/or **REDO** depends on **Buffer Manager**

- UNDO required if: uncommitted data can overwrite committed data (STEAL buffer management)
- REDO required if: transaction can commit before all its updates are on disk (NO FORCE buffer management)

The Log (cont.)

The following actions are recorded in the log:

- *if T_i writes an object, write a log record with:*
 - If UNDO required need “before image
 - IF REDO required need “after image”
- *T_i commits/aborts:* a log record indicating this action

Logging (cont.)

Write-Ahead Logging protocol

- Log record must go to disk before the changed page!
- All log records for a transaction (including its commit record) must be written to disk before the transaction is considered “Committed”

All logging and CC-related activities are handled transparently by the DBMS

(Review) Goal: The **ACID** properties

A tomicity: All actions in the transaction happen, or none happen

C onsistency: If each transaction is consistent, and the DB starts consistent, it ends up consistent

I solation: Execution of one transaction is isolated from that of other transactions

D urability: If a transaction commits, its effects persist

What happens if system **crashes** between
commit and flushing modified data to disk ?

D Durability - Recovering From a Crash

Three phases:

- Analysis: Scan the log (forward from the most recent *checkpoint*) to identify all transactions that were active at the time of the crash
- Redo: Redo updates as needed to ensure that all logged updates are in fact carried out and written to disk
- Undo: Undo writes of all transactions that were active at the crash, working backwards in the log

At the end – all committed updates and only those updates are reflected in the database

Some care must be taken to handle the case of a crash occurring during the recovery process!

Summary

Concurrency control and **recovery** are among the most important functions provided by a DBMS

Concurrency control is automatic

- System automatically inserts lock/unlock requests and schedules actions of different Xacts
- Property ensured: resulting execution is equivalent to executing the Xacts one after the other in some order

Write-ahead logging (WAL) and the recovery protocol are used to:

- 1.** undo the actions of aborted transactions, and
- 2.** restore the system to a consistent state after a crash

next: concurrency control in detail!